



Analyst Intern, Data Science & Solutions Project

William Rice

08/26/23

Introduction

The purpose of this project is to gauge your technical skills and problem solving ability by working through something science project. You will work your way through this R Markdown document, answering questions as you go along. Fill in your name to the "author" key in the YAML header. When you're finished with the document, come back and type your answers at the top. Please leave all your work below and have your answers where indicated below as well. Please note that we will so make it clear, concise and avoid long printouts. Feel free to add in as many new code chunks as you'd like.

Remember that we will be grading the quality of your code and visuals alongside the correctness of your answers. Please provide as much as possible (instead of base R and explicit loops.)

Note:

Throughout this document, any `season` column represents the year each season started. For example, the 2015 dataset is for the 2015-16 season. For most of the rest of the project, we will refer to a season by just this number (e.g. 2015) in the dataset as 2015-16).

Answers

Part 1

Question 1:

- 1st Team: 25.8 points per game
- 2nd Team: 23.1 points per game
- 3rd Team: 20.5 points per game
- All-Star: 21.6 points per game

Question 2: 4.7 Years

Question 3:

- Elite: 2 players.
- All-Star: 1 players.
- Starter: 10 players.
- Rotation: 8 players.
- Roster: 14 players.
- Out of League: 38 players.

Open Ended Modeling Question: Please show your work and leave all responses below in the document.

Part 2

Question 1: 28.9%

Question 2: Written question, put answer below in the document.

Question 3: Written question, put answer below in the document.

Setup and Data

```
library(tidyverse)
# Note, you will likely have to change these paths. If your data is in the same folder as this
# the paths will likely be fixed for you by deleting ../../Data/awards_project/ from each statement
awards <- read_csv("awards_data.csv")
player_data <- read_csv("player_stats.csv")
team_data <- read_csv("team_stats.csv")
rebounding_data <- read_csv("team_rebounding_data_22.csv")
```



Part 1 – Awards

In this section, you're going to work with data relating to player awards and statistics. You'll start with some data manipulation work towards building a model to predict broad levels of career success.

Question 1

QUESTION: What is the average number of points per game for players in the 2007-2021 seasons who won All NBA teams (**not** the All Defensive Teams), as well as for players who were in the All-Star Game (**not** the rookie all-star game)?

Here and for all future questions, feel free to add as many code chunks as you like. Do not worry about the order of the code chunks, we'll want to see your code.

```
#I am using separate code chunks to represent my work on data manipulation, then one code chunk for the final answer.
#I joined the awards and player statistics datasets together to have all relevant information so by the common season and nbapersonid features. Then, I created a points per game variable for the answers for question 1, by dividing total points that season by games.

library(dplyr)
part1df <- player_data %>% inner_join(awards,
  by=c('nbapersonid'='nbapersonid', 'season'='season'), relationship = 'many-to-many')
mutate(ppg=points/games)
```

#For each of the next 4 chunks, I created a smaller subset that only included players selected to the All NBA First Team. I then used summarize to average out ppg for each player and report it back. I did this for all 4 values identically.

```
#1st Team Work#
first_team_ppg <- part1df %>%
  filter(`All NBA First Team` == 1) %>%
  summarize(mean_ppg1 = mean(ppg))
first_team_ppg
```

```
## # A tibble: 1 × 1
##   mean_ppg1
##   <dbl>
## 1      25.9
```

```
#2nd Team Work#
second_team_ppg <- part1df %>%
  filter(`All NBA Second Team` == 1) %>%
  summarize(mean_ppg2 = mean(ppg))
second_team_ppg
```

```
## # A tibble: 1 × 1
##   mean_ppg2
##   <dbl>
## 1      23.1
```

```
#3rd Team Work#
third_team_ppg <- part1df %>%
  filter(`All NBA Third Team` == 1) %>%
  summarize(mean_ppg3 = mean(ppg))
third_team_ppg
```

```
## # A tibble: 1 × 1
##   mean_ppg3
##   <dbl>
## 1      20.5
```



```
#All-Star Work#
all_star_ppg <- part1df %>%
  filter(`all_star_game` == 'TRUE') %>%
  summarize(mean_ppg4 = mean(ppg))
all_star_ppg
```

```
## # A tibble: 1 × 1
##   mean_ppg4
##   <dbl>
## 1      21.6
```

#I made a graph on top of the 4 values shown just to visualize the differences in average ppg l-nba teams and all-star teams (and perhaps to do a little extra!). I used a dataframe to store values corresponding to the honors team, and also added custom colors for each one. Using gg features I was able to make a small plot for all four values, as shown below.#


```
honors_team <- c("First-Team All NBA", "Second-Team All NBA", "Third-Team All NBA", "All-Star")
avg_ppg <- c(sum(first_team_ppg), sum(second_team_ppg), sum(third_team_ppg), sum(all_star_ppg))

avg_ppg_plot_data <- data.frame(Honors = honors_team, PPG = avg_ppg)

plot_colors <- c('All-Star' = "blue", 'First-Team All NBA' = "orange", 'Second-Team All NBA' = "red", 'Third-Team All NBA' = "yellow")

ggplot(avg_ppg_plot_data, aes(x = Honors, y = PPG, fill = Honors)) +
  geom_bar(stat = "identity") +
  scale_fill_manual(values = plot_colors) +
  labs(title = "Average Points Per Game of Players Selected to Different Honors Teams", x = "Honors Teams", y = "Average PPG") +
  theme(axis.text.x = element_text(size = 8))
```

Average Points Per Game of Players Selected to Different Honors Teams



ANSWER 1:

1st Team: 25.8 points per game
 2nd Team: 23.1 points per game
 3rd Team: 20.5 points per game
 All-Star: 21.6 points per game



Question 2

QUESTION: What was the average number of years of experience in the league it takes for players to make their first or 3rd team)? Please limit your sample to players drafted in 2007 or later who did eventually go on to win at least one All-NBA award. Please provide an example:

- Luka Doncic is in the dataset as 2 years. He was drafted in 2018 and won his first All NBA award in 2019 (which was his 3rd season).
- LeBron James is not in this dataset, as he was drafted prior to 2007.
- Lu Dort is not in this dataset, as he has not received any All NBA honors.

```
#I took the large dataset that had combined player and awards statistics, and I filtered it to only
players drafted after 2006, and those who have been on any All-NBA team. I then created a new variable
called years_until_allNBA, which represents how long it took for each record to receive their All-NBA selection from the draft year
as people play in the same year they are drafted, and are given the award the end of the last season. This means that if I were to not add the plus one it would be a year behind. For example, if someone
named to an All-NBA team their rookie year, without the +1 it would show 0 years, which is incorrect. I added +1 to the data. Then, I grouped each player by their name, and filtered the data to only
the first season they received an award. This was the variable being asked for, so I then averaged the years_until_allNBA for all players and found the answer.
```

```
mean_years_experience <- part1df %>%
  filter(draftyear > 2006,
         `All NBA First Team` == 1 | `All NBA Second Team` == 1 | `All NBA Third Team` == 1)
mutate(years_until_allNBA = season - draftyear + 1) %>%
  group_by(player) %>%
  filter(season == min(season)) %>%
  ungroup() %>%
  summarize(mean_years = mean(years_until_allNBA))

mean_years_experience
```

```
## # A tibble: 1 × 1
##   mean_years
##       <dbl>
## 1         4.67
```

ANSWER 2:

4.7 Years

Data Cleaning Interlude

You're going to work to create a dataset with a "career outcome" for each player, representing the highest level of success achieved for **at least two** seasons *after his first four seasons in the league* (examples to follow below!). To do this, you need to define career outcomes. On a single season level, the outcomes are:

- Elite: A player is "Elite" in a season if he won any All NBA award (1st, 2nd, or 3rd team), MVP, or DPOY in that season.
- All-Star: A player is "All-Star" in a season if he was selected to be an All-Star that season.
- Starter: A player is a "Starter" in a season if he started in at least 41 games in the season OR if he played at least 900 minutes in the season.
- Rotation: A player is a "Rotation" player in a season if he played at least 1000 minutes in the season.
- Roster: A player is a "Roster" player in a season if he played at least 1 minute for an NBA team but did not meet the criteria for Starter, Rotation, or Elite.
- Out of the League: A player is "Out of the League" if he is not in the NBA in that season.

We need to make an adjustment for determining Starter/Rotation qualifications for a few seasons that didn't have 82 games. Specifically, if a player played 900 minutes in 2011, he **would** meet the rotation criteria because his minutes would be considered to be $900 * (82/66) = 1118$. Please use this math for both minutes and games started, so a player who started 38 games in 2020 would be considered to have started $38 * (82/72) = 43$ games, and thus would qualify for starting 41. Any answer assuming you round the multiplied values to the nearest whole number.

Note that on a season level, a player's outcome is the highest level of success he qualifies for in that season. Thus, if a player was both All-NBA 1st team and an All-Star last year, he would be considered to be "Elite" for the 2022 season. For a career outcome of All-Star if in the rest of his career he made one more All-Star game but no more All-NBA team appearances, and Shai has not yet played enough to have a career outcome.

Examples:



- A player who enters the league as a rookie and has season outcomes of Roster (1), Rotation (2), Rotation (3), and Out of the League (6+) would be considered “Out of the League,” because after his first four seasons, he only has a season in the league and does not qualify him for any success outcome.
- A player who enters the league as a rookie and has season outcomes of Roster (1), Rotation (2), Starter (3), Starter (6), All-Star (7), Elite (8), Starter (9) would be considered “All-Star,” because he had at least two seasons after his first season in production or higher.
- A player who enters the league as a rookie and has season outcomes of Roster (1), Rotation (2), Starter (3), Starter (6), Rotation (7), Rotation (8), Roster (9) would be considered a “Starter” because he has two seasons after his first season in production.

Question 3

QUESTION: There are 73 players in the `player_data` dataset who have 2010 listed as their draft year. How many of these players have each career outcome in each of the 6 buckets?

#Note: In the following r code chunk, I am very aware this is not concise. This took me a lot of time to write, and while I am most likely sure there were multiple ways to find this easier, this was the way I found to get the second highest season outcome when the frequency of the highest season outcome is 1. I combine them with the myriad of ifelse statements you see below. Again, I am aware this makes the code less concise and correctness, but overall I do believe my thought process of how to complete this task as well as my range of r knowledge in both tidyverse and base r was shown. Hopefully you understand my reasoning for this explanation but I hope this suffices! -Will Rice#



#I created a new function altogether that would account for those specific seasons needing to be added to the lockout and covid halts in play, respectively. I used if else so as to not change the results if the seasons were not changed in terms of max games played.#

```
calculate_mins <- function(mins, season) {
  ifelse(season == 2011, mins * (82/66), ifelse(season %in% c(2019, 2020), mins * (82/72), mins))
}
```

#Next, I once again joined the two datasets together for a fresh new dataset, so I can create a new question: Highest Outcome. I followed the specific criteria asked for in terms of different cases when function, then implemented my new minutes/games variable to alter those 3 seasons' results. I also rounded them as the first time I viewed the dataset it looked off as all of the outcomes were nice and rounded. Then, I filtered for the 2010 draft year, as that was part of the question included all player's first four years since being drafted as those did not count towards a career by player and season and made sure there were no duplicate years: there were players that the same season they had duplicated values that would count for them and in most cases help them get a similar playing time or better, which would not be fair, as that only occurred in one season so only one would show up. Then I grouped by player and their highest outcome, and found the highest outcome, which was ranked in the order I told R to rank using the arrange(match) function. Then, I used the slice function to find the highest outcome each player had during their 5th year. If a player's highest outcome had a frequency of 2, it was listed as their highest outcome. If their highest outcome had a frequency of 1, they were listed as out of the league for the logical reasons that the second highest outcome means the other years they could only be out of the league. Everyone else was NA. The process is below.#

```
career_outcomes <- player_data %>%
  left_join(awards, by = c("nba_personid", "season"), relationship = 'many-to-many') %>%
  mutate(HighestOutcome = case_when(
    `All NBA First Team` == 1 | `All NBA Second Team` == 1 | `All NBA Third Team` == 1 |
    `Most Valuable Player_rank` == 1 | `Defensive Player Of The Year_rank` == 1 ~ "Elite",
    all_star_game == 'TRUE' ~ "All-Star",
    games_start >= 41 | mins >= 2000 ~ "Starter",
    mins >= 1000 ~ "Rotation",
    mins >= 1 ~ "Roster",
    TRUE ~ "Out of the League"
  )) %>%
  mutate(mins = calculate_mins(mins, season),
         games_start = calculate_mins(games_start, season),
         mins = as.integer(round(mins)),
         games_start = as.integer(round(games_start))) %>%
  filter(season >= draftyear + 4, draftyear == 2010) %>%
  group_by(player, season) %>%
  filter(row_number() == 1) %>%
  ungroup() %>%
  group_by(player, HighestOutcome) %>%
  summarize(Frequency = n(), .groups = "keep") %>%
  arrange(match(HighestOutcome, c("Elite", "All-Star", "Starter", "Rotation", "Roster", "Out of the League"))) %>%
  group_by(player) %>%
  slice(1) %>%
  mutate(CareerOutcome = if_else(Frequency >= 2, HighestOutcome, NA),
         CareerOutcome = if_else(HighestOutcome == 'Roster' & Frequency == 1, 'Out of the League', HighestOutcome)) %>%
  group_by(player) %>%
  ungroup()
```

#This is where I was left frustrated and ended up happy with the outcome, though I realize I how often I'm creating new datasets. I created another dataset, identical to the previous one. I created the second_highest_outcome variable, it is mostly the same process. I created second_highest_outcome is the same as highest outcome except using some if_else statements and named differently for the second highest outcome. Next, I counted frequencies and listed importance of each outcome in order once more, except this time to find the second highest outcome for each player. I joined the two datasets of each player's highest and second highest outcome (although it only listed the highest if the second highest was not NA). Deselecting the frequencies then left me with a simple dataset of each player's outcomes.

```
single_season_outcomes <- player_data %>%
```



```

left_join(awards, by = c("nbapersonid", "season"), relationship = 'many-to-many') %>%
mutate(Elite = (`All NBA First Team` == 1 | `All NBA Second Team` == 1 | `All NBA Third Te
Player Of The Year_rk` == 1 | `Most Valuable Player_rk` == 1) + 0,
Elite = ifelse(is.na(Elite), 0, Elite)) %>%
mutate(AllStar = (all_star_game == 'TRUE' & Elite == 0) + 0,
AllStar = ifelse(is.na(AllStar), 0, AllStar)) %>%
mutate(Starter = ((games_start > 40 | mins > 1999) & Elite == 0 & AllStar == 0) + 0,
Starter = ifelse(is.na(Starter), 0, Starter)) %>%
mutate(Rotation = (mins > 999 & Elite == 0 & AllStar == 0 & Starter == 0) + 0,
Rotation = ifelse(is.na(Rotation), 0, Rotation)) %>%
mutate(Roster = (mins > 0 & Elite == 0 & AllStar == 0 & Starter == 0 & Rotation == 0) + 0,
Roster = ifelse(is.na(Roster), 0, Roster)) %>%
mutate(OutOfLeague = ifelse(Elite == 0 & AllStar == 0 & Starter == 0 & Rotation == 0 & Ros
mutate(mins = if_else(season == 2011, mins*(82/66), mins)) %>%
mutate(mins = if_else(season == 2019|season == 2020, mins*(82/72), mins)) %>%
mutate(games_start = if_else(season == 2011, games_start*(82/66), games_start)) %>%
mutate(games_start = if_else(season == 2019|season == 2020, games_start*(82/72), games_sta
mutate(mins = as.integer(round(mins))) %>%
mutate(games_start = as.integer(round(games_start))) %>%
select(nbapersonid, draftyear, player, season, Elite, AllStar, Starter, Rotation, Roster,
filter(draftyear == 2010, season > 2013) %>%
group_by(player, season) %>%
filter(row_number() == 1) %>%
ungroup() %>%
mutate(SecondHighestOutcome = if_else(Elite == 1, 'Elite', if_else(AllStar == 1, 'All-Star
1, 'Starter', if_else(Rotation == 1, 'Rotation', if_else(Roster == 1, 'Roster', 'Out of Leag
group_by(player, SecondHighestOutcome) %>%
summarize(Frequency = n(), .groups = "keep") %>%
arrange(match(SecondHighestOutcome, c("Elite", "All-Star", "Starter", "Rotation", "Roster"
e")))) %>%
group_by(player) %>%
slice(2)

```

```
second_highest_outcomes <- single_season_outcomes
```

```

total_2010_outcomes <- left_join(career_outcomes, second_highest_outcomes, by = "player") %>
select(-Frequency.x, -Frequency.y) %>%
mutate(CareerOutcome = if_else(is.na(CareerOutcome), SecondHighestOutcome, CareerOutcome))
select(-SecondHighestOutcome)

```

*#I replaced all NA values in the highest outcome with its second highest outcome, and then d
est outcome. Next, I used the summarize function to count how many of each career outcome th
d. I also noticed I only retained 44 values after the 4 year of playing. This meant that if
e 73, and only 44 had a distinct outcome, there were 29 who never played in their fifth year
lly qualify as out of league. I created a variable that counted 29 more out of leagues in th
reated one table 'outcome_counts' that showed the counts of all 6 different outcomes for the*

```

outcome_counts <- total_2010_outcomes %>%
group_by(CareerOutcome) %>%
summarize(count = n())

out_of_league_before_5th_year <- data.frame(CareerOutcome = "Out of League", count = 29)

outcome_counts <- bind_rows(outcome_counts, out_of_league_before_5th_year) %>%
group_by(CareerOutcome) %>%
summarize(count = sum(count))
outcome_counts

```

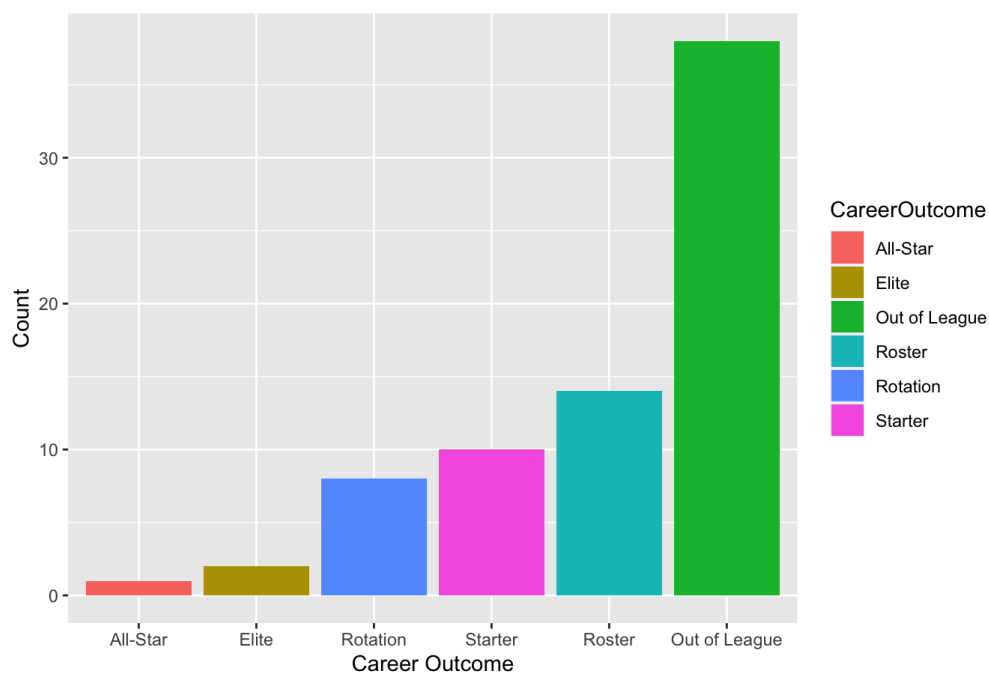


```
## # A tibble: 6 × 2
##   CareerOutcome count
##   <chr>          <dbl>
## 1 All-Star        1
## 2 Elite           2
## 3 Out of League  38
## 4 Roster         14
## 5 Rotation        8
## 6 Starter        10
```

#Next, I created again a ggplot visual to show counts by outcome, just to display how hard it is to thrive in the league, but thrive! I thought with some interesting colors generically given that it is a bit of a struggle to thrive in this otherwise code-heavy question due to my trouble with the second highest outcome.

```
ggplot(outcome_counts, aes(x = reorder(CareerOutcome, count), y = count, fill = CareerOutcome)) +
  geom_bar(stat = "identity") +
  labs(x = "Career Outcome", y = "Count", title = "Career Outcomes Across 2010 Draft Class")
```

Career Outcomes Across 2010 Draft Class



ANSWER 3:

Elite: 2 players.
 All-Star: 1 players.
 Starter: 10 players.
 Rotation: 8 players.
 Roster: 14 players.
 Out of League: 38 players.

Open Ended Modeling Question

In this question, you will work to build a model to predict a player's career outcome based on information up through career.

This question is intentionally left fairly open ended, but here are some notes and specifications.

1. We know modeling questions can take a long time, and that qualified candidates will have different levels of experience with modeling. Don't be discouraged. It's not our intention to make you spend excessive time here. If you get your ideas, think you could do better by spending a lot more time, you can just write a bit about your ideas for future improvement. Further, we're more interested in your thought process and critical thinking than we are in specific modeling techniques. Feature engineering is more important than using fancy mathematical machinery, and a successful candidate could use a simple model.
2. You may use any data provided in this project, but please do not bring in any external sources of data. Note that the data provided goes back to 2007, All NBA and All Rookie team voting is only included back to 2011.



3. A player needs to complete at least three additional seasons after their first four to be considered as having a career outcome in our dataset. (We are using 3+ instead of 2+ just to give each player a little more time to accumulate high level stats in his career). Because the dataset in this project ends in 2021, this means that a player would need to have had '20, and '19 seasons after his first four years, and thus his first four years would have been '18, '17, '16, and '15. **your training data to players who were drafted in or before the 2015 season.** Karl-Anthony Towns was the first player to do this.
4. Once you build your model, predict on all players who were drafted in 2018-2021 (They have between 1 and 4 seasons and have not yet started accumulating seasons that inform their career outcome).
5. You can predict a single career outcome for each player, but it's better if you can predict the probability that each player falls into each outcome bucket.
6. Include, as part of your answer:
 - A brief written overview of how your model works, targeted towards a decision maker in the front office without a data science background.
 - What you view as the strengths and weaknesses of your model.
 - How you'd address the weaknesses if you had more time and/or more data.
 - A ggplot or ggplotly visualization highlighting some part of your modeling process, the model itself, or your results.
 - Your predictions for Shai Gilgeous-Alexander, Zion Williamson, James Wiseman, and Josh Giddey.
 - (Bonus!) An html table (for example, see the package `reactable`) containing all predictions for the players drafted in or before the 2015 season.

#Again, I created a master dataset that joined together awards and player data, mutating any NA values to a 0. Then, I went back and remedied a few choice columns: draftpick was given an NA again if the player was all nba and all rookie voting ranks before 2011, as there was no data. Then, I created career outcome probabilities and rounded them to one decimal point, while I deselected some of the stats I did not believe had a significant correlation to a potential career outcome, or that were already a part of the model.

#MODEL DATASET#

```
model_data <- left_join(player_data, awards, by = c("nbapersonid", "season"), relationship = "one-to-many")
mutate_all(~ replace(., is.na(.), 0)) %>%
mutate(draftpick = if_else(draftpick == 0, NA, draftpick)) %>%
mutate(all_nba_points_rk = if_else(all_nba_points_rk == 0 & season < 2011, NA, all_nba_points_rk)) %>%
mutate(all_rookie_points_rk = if_else(all_rookie_points_rk == 0 & season < 2011, NA, all_rookie_points_rk)) %>%
filter(draftyear >= 2003) %>%
mutate(ppg = (points/games)) %>%
mutate(rpg = (tot_reb/games)) %>%
mutate(apg = (ast/games)) %>%
mutate(bpg = (blocks/games)) %>%
mutate(spg = (steals/games)) %>%
mutate_at(vars(apg, spg, bpg, ppg, rpg), ~ round(., 1)) %>%
select(-draftpick, -nbateamid, -'Player Of The Month', -'Rookie Of The Month', -'Player Of The Year',
points_rk, -all_rookie_points_rk, -allstar_rk)
```

model_data

A tibble: 7,279 × 67

	nbapersonid	player	draftyear	season	team	games	games_start	mins	fgm	fga
##	<dbl>	<chr>	<dbl>	<dbl>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
##	1	2585 Zaza ...	2003	2007	ATL	62	5	944	107	245
##	2	200780 Solom...	2006	2007	ATL	35	0	145	12	30
##	3	2746 Josh ...	2004	2007	ATL	81	81	2873	518	1133
##	4	201151 Acie ...	2007	2007	ATL	56	6	865	95	237
##	5	101136 Salim...	2005	2007	ATL	35	0	402	65	180
##	6	2735 Josh ...	2004	2007	ATL	76	0	2274	327	573
##	7	200978 Jerem...	2006	2007	ATL	19	0	88	13	31
##	8	201143 Al Ho...	2007	2007	ATL	81	77	2540	333	668
##	9	101107 Marvi...	2005	2007	ATL	80	80	2765	424	918
##	10	200749 Sheld...	2006	2007	ATL	36	0	414	37	100

i 7,269 more rows

i 57 more variables: fgp <dbl>, fgm3 <dbl>, fga3 <dbl>, fgp3 <dbl>, fgm2 <dbl>, fga2 <dbl>, fgp2 <dbl>, efg <dbl>, ftm <dbl>, fta <dbl>, ftp <dbl>, off_reb <dbl>, def_reb <dbl>, tot_reb <dbl>, ast <dbl>, steals <dbl>, blocks <dbl>, tov <dbl>, tot_fouls <dbl>, points <dbl>, PER <dbl>, FTr <dbl>, off_reb_pct <dbl>, def_reb_pct <dbl>, tot_reb_pct <dbl>, ast_pct <dbl>, stl_pct <dbl>, blk_pct <dbl>, ...



```

calculate_mins <- function(mins, season) {
  ifelse(season == 2011, mins * (82/66), ifelse(season %in% c(2019, 2020), mins * (82/72), m
})

#Next, I made my training dataset. I took the model dataset and filtered out only players dr
2015. This was imperative to me as to have all of the first four years of data on every play
as well as the years beyond that which made to help generate a distinct outcome. For the nex
I reproduced finding the highest career outcome after doing some duplicate codes to ensure t
eft with the next highest seasonal outcome.

#TRAINING DATASET#

training_set <- model_data %>%
  filter(draftyear < 2016 & draftyear > 2006) %>%
  mutate(SeasonOutcome = case_when(
    `All NBA First Team` == 1 | `All NBA Second Team` == 1 | `All NBA Third Team` == 1 |
    `Most Valuable Player_rk` == 1 | `Defensive Player Of The Year_rk` == 1 ~ "Elite",
    all_star_game == 1 ~ "All-Star",
    games_start >= 41 | mins >= 2000 ~ "Starter",
    mins >= 1000 ~ "Rotation",
    mins >= 1 ~ "Roster",
    TRUE ~ "Out of the League"
  )) %>%
  mutate(mins = calculate_mins(mins, season),
    games_start = calculate_mins(games_start, season),
    mins = as.integer(round(mins)),
    games_start = as.integer(round(games_start))) %>%
  filter(season >= draftyear + 4) %>%
  group_by(player, season) %>%
  filter(row_number() == 1) %>%
  ungroup() %>%
  group_by(player, SeasonOutcome) %>%
  summarize(Frequency = n(), .groups = "keep") %>%
  arrange(match(SeasonOutcome, c("Elite", "All-Star", "Starter", "Rotation", "Roster", "Out
group_by(player) %>%
  slice(1) %>%
  mutate(CareerOutcome = if_else(Frequency >= 3, SeasonOutcome, NA),
    CareerOutcome = if_else(SeasonOutcome == 'Roster' & Frequency == 1, 'Out of League'
group_by(player) %>%
  ungroup()

training_data_2 <- model_data %>%
  filter(draftyear < 2016 & draftyear > 2006) %>%
  mutate(SecondHighestSeasonOutcome = case_when(
    `All NBA First Team` == 1 | `All NBA Second Team` == 1 | `All NBA Third Team` == 1 |
    `Most Valuable Player_rk` == 1 | `Defensive Player Of The Year_rk` == 1 ~ "Elite",
    all_star_game == 1 ~ "All-Star",
    games_start >= 41 | mins >= 2000 ~ "Starter",
    mins >= 1000 ~ "Rotation",
    mins >= 1 ~ "Roster",
    TRUE ~ "Out of the League"
  )) %>%
  mutate(mins = calculate_mins(mins, season),
    games_start = calculate_mins(games_start, season),
    mins = as.integer(round(mins)),
    games_start = as.integer(round(games_start))) %>%
  filter(season >= draftyear + 4) %>%
  group_by(player, season) %>%
  filter(row_number() == 1) %>%
  ungroup() %>%
  group_by(player, SecondHighestSeasonOutcome) %>%
  summarize(Frequency = n(), .groups = "keep") %>%
  arrange(match(SecondHighestSeasonOutcome, c("Elite", "All-Star", "Starter", "Rotation", "R
eague"))) %>%
  group_by(player) %>%
  slice(2)

```



```
training_data_third_highest <- model_data %>%
  filter(draftyear < 2016 & draftyear > 2006) %>%
  mutate(ThirdHighestSeasonOutcome = case_when(
    `All NBA First Team` == 1 | `All NBA Second Team` == 1 | `All NBA Third Team` == 1 |
    `Most Valuable Player_rk` == 1 | `Defensive Player Of The Year_rk` == 1 ~ "Elite",
    all_star_game == 1 ~ "All-Star",
    games_start >= 41 | mins >= 2000 ~ "Starter",
    mins >= 1000 ~ "Rotation",
    mins >= 1 ~ "Roster",
    TRUE ~ "Out of the League"
  )) %>%
  mutate(mins = calculate_mins(mins, season),
    games_start = calculate_mins(games_start, season),
    mins = as.integer(round(mins)),
    games_start = as.integer(round(games_start))) %>%
  filter(season >= draftyear + 4) %>%
  group_by(player, season) %>%
  filter(row_number() == 1) %>%
  ungroup() %>%
  group_by(player, ThirdHighestSeasonOutcome) %>%
  summarize(Frequency = n(), .groups = "keep") %>%
  arrange(match(ThirdHighestSeasonOutcome, c("Elite", "All-Star", "Starter", "Rotation", "Ro
ague"))) %>%
  group_by(player) %>%
  slice(3)
```

```
View(training_data_third_highest)
```

```
total_training_outcomes_draft <- left_join(training_set, training_data_2, by = "player") %>%
  mutate(CareerOutcome = if_else(Frequency.x == 1, if_else(Frequency.y <= 1, NA, SecondHighe
eerOutcome)) %>%
  mutate(CareerOutcome = if_else(Frequency.x == 2, SecondHighestSeasonOutcome, CareerOutcome
  mutate(CareerOutcome = if_else(is.na(CareerOutcome) & is.na(SecondHighestSeasonOutcome), "
rOutcome))
View(total_training_outcomes_draft)
```

```
total_training_outcomes <- left_join(total_training_outcomes_draft, training_data_third_high
>%
  select(-SecondHighestSeasonOutcome, -Frequency.y) %>%
  mutate(CareerOutcome = if_else(is.na(CareerOutcome) & is.na(ThirdHighestSeasonOutcome), "O
Outcome)) %>%
  mutate(CareerOutcome = if_else(is.na(CareerOutcome), ThirdHighestSeasonOutcome, CareerOutc
View(total_training_outcomes)
```

*#I then joined to gether the dataset of all players drafted from 2007 and 2015 with the play
er outcomes, and then filtered all players who did not play long enough to 'Out of League.'
for my training data.*

```
training_data_3 <- model_data %>%
  filter(draftyear < 2016 & draftyear > 2006)

training_dataset_main <- left_join(training_data_3, total_training_outcomes, by = "player")
  mutate(CareerOutcome = if_else(is.na(CareerOutcome) & is.na(SeasonOutcome), 'Out of League
>%
  select(-SeasonOutcome)
View(training_dataset_main)
```

*#I then took their first four years of data as the training data. This is because we are pre
ata's first one-four years of data to make a prediction. We have career outcomes for each of
e can also aggregate their first four years to build a model that can help us evaluate the t*



ar trends in variables and data. I took the next massive few chunks of code to create aggreg the first four years of each players career. I believed having one record for each player in e to fit into a model better, and it just took a bit of time to write down every stat so as rectly.

```
first_4_years <- training_dataset_main %>%
  filter(season <= draftyear + 3) %>%
  group_by(player, season) %>%
  filter(row_number() == 1) %>%
  ungroup() %>%
  mutate(MVP_Award = if_else(`Most Valuable Player_rk` == 1, 1, 0)) %>%
  mutate(MVP_Candidate = if_else(`Most Valuable Player_rk` > 1 & `Most Valuable Player_rk` <
  mutate(MIP_Award = if_else(`Most Improved Player_rk` == 1, 1, 0)) %>%
  mutate(MIP_Candidate = if_else(`Most Improved Player_rk` > 1 & `Most Improved Player_rk` <
  mutate(DPOY_Award = if_else(`Defensive Player Of The Year_rk` == 1, 1, 0)) %>%
  mutate(DPOY_Candidate = if_else(`Defensive Player Of The Year_rk` > 1 & `Defensive Player
1, 0)) %>%
  mutate(ROY_Award = if_else(`Rookie Of The Year_rk` == 1, 1, 0)) %>%
  mutate(ROY_Candidate = if_else(`Rookie Of The Year_rk` > 1 & `Rookie Of The Year_rk` < 10,
  mutate(Sixth_Man_Award = if_else(`Sixth Man Of The Year_rk` == 1, 1, 0)) %>%
  mutate(Sixth_Man_Candidate = if_else(`Sixth Man Of The Year_rk` > 1 & `Sixth Man Of The Ye

first_4_years_training <- first_4_years %>%
  group_by(player) %>%
  summarize(
    all_star_games = sum(all_star_game),
    all_nba_defensive_first_team_selections = sum(`All NBA Defensive First Team`),
    all_nba_defensive_second_team_selections = sum(`All NBA Defensive Second Team`),
    all_nba_first_team_selections = sum(`All NBA First Team`),
    all_nba_second_team_selections = sum(`All NBA Second Team`),
    all_nba_third_team_selections = sum(`All NBA Third Team`),
    all_rookie_first_team_selections = sum(`All Rookie First Team`),
    all_rookie_second_team_selections = sum(`All Rookie Second Team`),
    MVP_count = sum(`MVP_Award`),
    MIP_count = sum(`MIP_Award`),
    DPOY_count = sum(`DPOY_Award`),
    ROY_count = sum(`ROY_Award`),
    sixth_man_count = sum(`Sixth_Man_Award`),
    MVP_candidate_count = sum(`MVP_Candidate`),
    MIP_candidate_count = sum(`MIP_Candidate`),
    DPOY_candidate_count = sum(`DPOY_Candidate`),
    ROY_candidate_count = sum(`ROY_Candidate`),
    sixth_man_candidate_count = sum(`Sixth_Man_Candidate`),
    finals_mvp = sum(`Bill Russell NBA Finals MVP`),
    total_points = sum(points),
    total_assists = sum(points),
    total_rebounds = sum(tot_reb),
    total_steals = sum(steals),
    total_blocks = sum(blocks),
    total_turnovers = sum(tov),
    total_field_goals = sum(fgm),
    total_field_goals_attempted = sum(fga),
    total_two_pointers = sum(fgm2),
    total_two_pointers_attempted = sum(fga2),
    total_three_pointers = sum(fgm3),
    total_three_pointers_attempted = sum(fga3),
    total_free_throws = sum(ftm),
    total_free_throws_attempted = sum(fta),
    total_games_started = sum(games_start),
    total_minutes = sum(mins),
    total_defensive_win_shares = sum(DWS),
    total_offensive_win_shares = sum(OWS),
    total_win_shares = sum(WS),
    total_defensive_box_plus_minus = sum(DBPM),
    total_offensive_box_plus_minus = sum(OBPM),
    total_box_plus_minus = sum(BPM),
    total_VORP = sum(VORP),
    total_games = sum(games),
    total_PER = sum(PER),
```



```

total_usg = sum(usg)

)

first_4_years_aggregated <- first_4_years_training %>%
  mutate(ppg = (total_points/total_games)) %>%
  mutate(apg = (total_assists/total_games)) %>%
  mutate(rpg = (total_rebounds/total_games)) %>%
  mutate(bpg = (total_blocks/total_games)) %>%
  mutate(spg = (total_steals/total_games)) %>%
  mutate(topg = (total_turnovers/total_games)) %>%
  mutate(fgpg = (total_field_goals/total_games)) %>%
  mutate(fgapg = (total_field_goals_attempted/total_games)) %>%
  mutate(fgpercentage = (fgpg/fgapg)) %>%
  mutate(fg2pg = (total_two_pointers/total_games)) %>%
  mutate(fg2apg = (total_two_pointers_attempted/total_games)) %>%
  mutate(fg2percentage = (fg2pg/fg2apg)) %>%
  mutate(fg3pg = (total_three_pointers/total_games)) %>%
  mutate(fg3apg = (total_three_pointers_attempted/total_games)) %>%
  mutate(fg3percentage = (fg3pg/fg3apg)) %>%
  mutate(ftpg = (total_free_throws/total_games)) %>%
  mutate(ftapg = (total_free_throws_attempted/total_games)) %>%
  mutate(ftpercentage = (ftpg/ftapg)) %>%
  mutate(mpg = (total_minutes/total_games)) %>%
  mutate(average_DWS = (total_defensive_win_shares/4)) %>%
  mutate(average_OWS = (total_offensive_win_shares/4)) %>%
  mutate(average_WS = (total_win_shares/4)) %>%
  mutate(average_DBPM = (total_defensive_box_plus_minus/4)) %>%
  mutate(average_OBPM = (total_offensive_box_plus_minus/4)) %>%
  mutate(average_BPM = (total_box_plus_minus/4)) %>%
  mutate(average_VORP = (total_VORP/4)) %>%
  mutate(average_games = (total_games/4)) %>%
  mutate(average_games_started = (total_games_started/4)) %>%
  mutate(average_usg = (total_usg/4)) %>%
  mutate(average_PER = (total_PER/4)) %>%
  select(-total_points, -total_assists, -total_rebounds, -total_blocks, -total_steals, -total_field_goals, -total_field_goals_attempted, -total_two_pointers, -total_two_pointers_attempted, -total_three_pointers_attempted, -total_free_throws, -total_free_throws_attempted, -total_minutes, -total_games, -total_VORP, -total_defensive_win_shares, -total_offensive_win_shares, -total_defensive_box_plus_minus, -total_offensive_box_plus_minus, -total_box_plus_minus)
  mutate_at(vars(ppg, apg, rpg, bpg, spg, topg, fgpg, fgapg, fg2pg, fg2apg, fg3pg, fg3apg, fgpercentage, fg2percentage, fg3percentage, ftpercentage, average_BPM, average_DWS, average_OWS, average_WS, average_VORP), ~ round(., 3))

und(., 1)) %>%
  mutate_at(vars(average_games_started, average_games), ~ round(., 0)) %>%
  mutate_at(vars(fgpercentage, fg2percentage, fg3percentage, ftpercentage, average_BPM, average_DWS, average_OWS, average_WS, average_VORP), ~ round(., 3))

#I next joined together the two datasets that had all of the aggregate stats with other stats like the name of the player, draft year, actual career outcome, etc. I created a small change in the range in the first four years they played in stead of four separate identical records except for each player.

model_training_dataset <- left_join(first_4_years_aggregated, first_4_years %>% select(nba_player_id, draftyear, CareerOutcome), by = "player") %>%
  group_by(player) %>%
  mutate(years = paste(min(season), max(season), sep = "-"), .groups = "drop") %>%
  select(-season, -.groups) %>%
  ungroup() %>%
  distinct(player, .keep_all = TRUE)
View(model_training_dataset)

#I did the same thing as the training dataset, instead using players drafted in or after 2017 training dataset. I then created years of experience to help with some of the averages per year. A player in the training had 4 years of data, most of these players have varying amounts between their individual experience amount and used that variable to divide by specific statistics.

#TESTING DATASET#

```



```

testing_dataset <- model_data %>%
  filter(draftyear >= 2018) %>%
  group_by(player, season) %>%
  filter(row_number() == 1) %>%
  ungroup() %>%
  mutate(MVP_Award = if_else(`Most Valuable Player_rk` == 1, 1, 0)) %>%
  mutate(MVP_Candidate = if_else(`Most Valuable Player_rk` > 1 & `Most Valuable Player_rk` <
  mutate(MIP_Award = if_else(`Most Improved Player_rk` == 1, 1, 0)) %>%
  mutate(MIP_Candidate = if_else(`Most Improved Player_rk` > 1 & `Most Improved Player_rk` <
  mutate(DPOY_Award = if_else(`Defensive Player Of The Year_rk` == 1, 1, 0)) %>%
  mutate(DPOY_Candidate = if_else(`Defensive Player Of The Year_rk` > 1 & `Defensive Player
1, 0)) %>%
  mutate(ROY_Award = if_else(`Rookie Of The Year_rk` == 1, 1, 0)) %>%
  mutate(ROY_Candidate = if_else(`Rookie Of The Year_rk` > 1 & `Rookie Of The Year_rk` < 10,
  mutate(Sixth_Man_Award = if_else(`Sixth Man Of The Year_rk` == 1, 1, 0)) %>%
  mutate(Sixth_Man_Candidate = if_else(`Sixth Man Of The Year_rk` > 1 & `Sixth Man Of The Ye
>%
  mutate(Years_Experience = season - draftyear + 1)

View(testing_dataset)

testing_dataset_summed_up <- testing_dataset %>%
  group_by(player) %>%
  summarize(
    all_star_games = sum(all_star_game),
    all_nba_defensive_first_team_selections = sum(`All NBA Defensive First Team`),
    all_nba_defensive_second_team_selections = sum(`All NBA Defensive Second Team`),
    all_nba_first_team_selections = sum(`All NBA First Team`),
    all_nba_second_team_selections = sum(`All NBA Second Team`),
    all_nba_third_team_selections = sum(`All NBA Third Team`),
    all_rookie_first_team_selections = sum(`All Rookie First Team`),
    all_rookie_second_team_selections = sum(`All Rookie Second Team`),
    MVP_count = sum(`MVP_Award`),
    MIP_count = sum(`MIP_Award`),
    DPOY_count = sum(`DPOY_Award`),
    ROY_count = sum(`ROY_Award`),
    sixth_man_count = sum(`Sixth_Man_Award`),
    MVP_candidate_count = sum(`MVP_Candidate`),
    MIP_candidate_count = sum(`MIP_Candidate`),
    DPOY_candidate_count = sum(`DPOY_Candidate`),
    ROY_candidate_count = sum(`ROY_Candidate`),
    sixth_man_candidate_count = sum(`Sixth_Man_Candidate`),
    finals_mvp = sum(`Bill Russell NBA Finals MVP`),
    total_points = sum(points),
    total_assists = sum(points),
    total_rebounds = sum(tot_reb),
    total_steals = sum(steals),
    total_blocks = sum(blocks),
    total_turnovers = sum(tov),
    total_field_goals = sum(fgm),
    total_field_goals_attempted = sum(fga),
    total_two_pointers = sum(fgm2),
    total_two_pointers_attempted = sum(fga2),
    total_three_pointers = sum(fgm3),
    total_three_pointers_attempted = sum(fga3),
    total_free_throws = sum(ftm),
    total_free_throws_attempted = sum(fta),
    total_games_started = sum(games_start),
    total_minutes = sum(mins),
    total_defensive_win_shares = sum(DWS),
    total_offensive_win_shares = sum(OWS),
    total_win_shares = sum(WS),
    total_defensive_box_plus_minus = sum(DBPM),
    total_offensive_box_plus_minus = sum(OBPM),
    total_box_plus_minus = sum(BPM),
    total_VORP = sum(VORP),
    total_games = sum(games),
    total_PER = sum(PER),
    total_usg = sum(usg),

```



```

years_experience = max(Years_Experience)

)

View(testing_dataset_summed_up)

testing_dataset_aggregated <- testing_dataset_summed_up %>%
  mutate(ppg = (total_points/total_games)) %>%
  mutate(apg = (total_assists/total_games)) %>%
  mutate(rpg = (total_rebounds/total_games)) %>%
  mutate(bpg = (total_blocks/total_games)) %>%
  mutate(spg = (total_steals/total_games)) %>%
  mutate(topg = (total_turnovers/total_games)) %>%
  mutate(fgpg = (total_field_goals/total_games)) %>%
  mutate(fgapg = (total_field_goals_attempted/total_games)) %>%
  mutate(fgpercentage = (fgpg/fgapg)) %>%
  mutate(fg2pg = (total_two_pointers/total_games)) %>%
  mutate(fg2apg = (total_two_pointers_attempted/total_games)) %>%
  mutate(fg2percentage = (fg2pg/fg2apg)) %>%
  mutate(fg3pg = (total_three_pointers/total_games)) %>%
  mutate(fg3apg = (total_three_pointers_attempted/total_games)) %>%
  mutate(fg3percentage = (fg3pg/fg3apg)) %>%
  mutate(ftpg = (total_free_throws/total_games)) %>%
  mutate(ftapg = (total_free_throws_attempted/total_games)) %>%
  mutate(ftpercentage = (ftpg/ftapg)) %>%
  mutate(mpg = (total_minutes/total_games)) %>%
  mutate(average_DWS = (total_defensive_win_shares/years_experience)) %>%
  mutate(average_OWS = (total_offensive_win_shares/years_experience)) %>%
  mutate(average_WS = (total_win_shares/years_experience)) %>%
  mutate(average_DBPM = (total_defensive_box_plus_minus/years_experience)) %>%
  mutate(average_OBPM = (total_offensive_box_plus_minus/years_experience)) %>%
  mutate(average_BPM = (total_box_plus_minus/years_experience)) %>%
  mutate(average_VORP = (total_VORP/years_experience)) %>%
  mutate(average_games = (total_games/years_experience)) %>%
  mutate(average_games_started = (total_games_started/years_experience)) %>%
  mutate(average_usg = (total_usg/years_experience)) %>%
  mutate(average_PER = (total_PER/years_experience)) %>%
  select(-total_points, -total_assists, -total_rebounds, -total_blocks, -total_steals, -total_field_goals, -total_field_goals_attempted, -total_two_pointers, -total_two_pointers_attempted, -total_three_pointers_attempted, -total_free_throws, -total_free_throws_attempted, -total_minutes, -total_games, -total_VORP, -total_defensive_win_shares, -total_offensive_win_shares, -total_defensive_box_plus_minus, -total_offensive_box_plus_minus, -total_box_plus_minus)
  mutate_at(vars(ppg, apg, rpg, bpg, spg, topg, fgpg, fgapg, fg2pg, fg2apg, fg3pg, fg3apg, fg3percentage, ftpg, ftapg, ftpercentage, mpg, average_games_started, average_games), ~ round(., 0)) %>%
  mutate_at(vars(fgpercentage, fg2percentage, fg3percentage, ftpercentage, average_BPM, average_DWS, average_OWS, average_WS, average_VORP), ~ round(., 3))

View(testing_dataset_aggregated)

```

#I created a formula that displayed CareerOutcome as the target variable, with almost every variable that had been aggregated over time, including sums of awards within the dataset. I then installed the random forest package and ran a random forest model on the training data, then tested it on the testing dataset. I changed some NA values to 0 instead of null).

#MODEL TRAINING, PREDICTIONS, AND EVALUATION#

```

Model_Formula <- CareerOutcome ~ all_star_games + all_nba_defensive_first_team_selections +
  all_nba_defensive_second_team_selections + all_nba_first_team_selections + all_nba_second_team_selections + all_nba_rookie_first_team_selections + all_nba_rookie_second_team_selections + MVP_count + MIP_count + DPOY_count + sixth_man_count + MVP_candidate_count + MIP_candidate_count + ROY_candidate_count + sixth_man_candidate_count + finals_mvp + average_PER + average_usg + ppg + apg + rpg + fgpg + fgapg + fgpercentage + fg2pg + fg2apg + fg2percentage + fg3pg + fg3apg + fg3percentage + ftpg + ftapg + ftpercentage + mpg + average_games + average_games_started + average_VORP + average_OWS + average_WS + average_DBPM + average_OBPM + average_BPM

```



```
#Installing Random Forest Model and setting seed for reproducibility#  
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:dplyr':  
##  
## combine
```

```
## The following object is masked from 'package:ggplot2':  
##  
## margin
```

```
set.seed(123)
```

```
#Ensuring All Null Values are 0.00%#
```

```
model_training_dataset$CareerOutcome <- factor(model_training_dataset$CareerOutcome)  
levels(model_training_dataset$CareerOutcome)
```

```
## [1] "All-Star"      "Elite"          "Out of League" "Roster"  
## [5] "Rotation"      "Starter"
```




```
View(model_training_dataset$CareerOutcome)

model_training_dataset <- model_training_dataset %>%
  mutate(ftpercentage = if_else(is.na(ftpercentage), 0, ftpercentage))

testing_dataset_aggregated <- testing_dataset_aggregated %>%
  mutate(ftpercentage = if_else(is.na(ftpercentage), 0, ftpercentage))

model_training_dataset <- model_training_dataset %>%
  mutate(fg3percentage = if_else(is.na(fg3percentage), 0, fg3percentage))

testing_dataset_aggregated <- testing_dataset_aggregated %>%
  mutate(fg3percentage = if_else(is.na(fg3percentage), 0, fg3percentage))

model_training_dataset <- model_training_dataset %>%
  mutate(fg2percentage = if_else(is.na(fg2percentage), 0, fg2percentage))

testing_dataset_aggregated <- testing_dataset_aggregated %>%
  mutate(fg2percentage = if_else(is.na(fg2percentage), 0, fg2percentage))

model_training_dataset <- model_training_dataset %>%
  mutate(fgpercentage = if_else(is.na(fgpercentage), 0, fgpercentage))

testing_dataset_aggregated <- testing_dataset_aggregated %>%
  mutate(fgpercentage = if_else(is.na(fgpercentage), 0, fgpercentage))

#Back to random forests models#

rf_model <- randomForest(Model_Formula, data = model_training_dataset)

testing_predictions <- predict(rf_model, newdata = testing_dataset_aggregated)

outcome_probs <- predict(rf_model, newdata = testing_dataset_aggregated, type = "prob")
View(outcome_probs)

outcome_predictions <- cbind(testing_dataset_aggregated[c("player")], testing_predictions)
View(outcome_predictions)

probability_predictions <- cbind(testing_dataset_aggregated[c("player")], outcome_probs)
View(probability_predictions)

#Below is my dataset configured in a later chunk that shows a player, his model-predicted ou
ies of landing in each of the six outcomes. The HTML table shows this too!#

total_predictions <- left_join(outcome_predictions, probability_predictions, by = "player")
  rename('Model Predictions' = 'testing_predictions')
View(total_predictions)

#Above is the evaluation and prediction outcomes of the testing dataset. One dataset has the
for each player, while the other specifically shows the outcomes available as a probability
w is my full description for the model, strengths and weaknesses, and potential solutions to
n the constraints.
```

BRIEF OVERVIEW OF MODEL AND HOW IT WORKS, STRENGTHS AND WEAKNESSES HOW TO ADDRESS WEAKNESSES WITH MORE TIME/KNOWLEDGE/EXPERIENCE

My model tries to follow the most basic guidelines of the project, and make the most out of the parameters given using a logical approach. First, I went and followed the steps given pretty closely: I took the data we have on individual players and combined them with the data on individual player awards, all over the span of 2007-2021. I followed the approach that we use as our 'training data', so as to give the model data to understand the relationships and variables in order to predict future outcomes more precisely. I took players in the 2007-2015 draft classes. That way, I have careers worth of data on every player looked at, as our dataset only went back to 2007, and using players before that would result in less preciseness. I used the guidelines for the season outcome classification as only deemed to have a distinct career outcome if he attains that level or higher on 3 different campaigns. I factored the dataset with issuing a career outcome for each player from the 2007-2015 class as using averages of all of their first



using four different records so as to use up less rows and cause less confusion for the computer and model. After this to data and concocting the testing set of players from 2018-2021 with the same statistical variables (as we must use the model is predicting on the right things), it was time for me to pick, train, predict, and evaluate with my model. I know variable, Career Outcome, is a categorical variable (non-numeric, one of few options, and if it is not the first five of six), but also non-binary, meaning there is more than two options. This leaves little wiggle room into which type of model also where my lack of expert statistical knowledge comes into play. The only few options I had thought of were random multinomial regressions, and I chose random forests as I had recently used this model to predict on another dataset and learning it is quite accurate. Random forests are very complex in their inner-workings and how they come to the conclusion process behind them and simplified version is easy to understand. A random forest takes a lot of decision trees, or simply takes a few variables and follows certain guidelines, made up of some of the different data points given and spliced in terms of patterns, order, etc. The random forest will look at the totality of the answers given from the decision trees, then the decision trees have to say and use the combined knowledge from them to give its output and predictions. This is how trees create combinations of data that end up being more accurate in its predictions than others; while the random forest it does not just listen to one decision tree, it listens to what feels to be the grand consensus of feature importance, and accuracy from all of the decision trees.

My model does have flaws (starting with how Josh Giddey is only considered a roster outcome!) In all seriousness, though constraints hindered me from displaying the most accurate, perfected model in terms of features, model type, and of course. First, I am only of certain level of knowledge of modeling, especially in the advanced multiclassification field. I am a precise modeling that may have been very accurate, but I also felt confident that like mentioned earlier in the guidelines would work. One weakness is potentially my feature choice. First, I could quickly go back and change variables and rather deliver my thought process with the original idea I had, which is the model who's predictions you see. I thought statistics and total award counts would suffice for a small-working model on short-time, as I did not want to pick any point of view about what I believe to be important in basketball. I let the model choose that first. I should have, however, summary statistics and looking for possible correlation patterns, or more importantly issues with multicollinearity. I did offensive, defensive, and total win shares, among other stats like this, and removing them may be beneficial to just go over. I also could have possibly removed or tweaked other statistics that may have been slightly multicollinear or repeated weights on certain statistics instead of using binary variables for some of the awards. It seems players with high-stat awards found success in career outcomes, while the model did not value some of the more important or telling awards. My last weakness I will mention in this report: team statistics. I did not include team statistics due to time constraints to figure out a perfect way to include them. Joining the datasets by player and year would have been no issue, I just was in time to develop and find something that could have helped offset high-volume and high-statistics players. For example, players who have massive point/assist/rebound totals, possibly due to being the focal point of an offense, or the best player on the team, usage rate may not have helped all that much. Another note on this is that potentially a few more defensive statistics on offense in these player statistics could have been helpful for players who do not score at will, and find their skills hidden. Lastly, back to team statistics, having team stats in relation to a player would have been interesting, as well as team time a player played there. It would be difficult with the flow of player across teams, but if possible, the model could use more contextual information regarding where a team was overall while a player was generating the stats seen in the report.

###VISUALS FOR MODELING QUESTION###

#Predicted Outcome Counts Graph#

```
desired_order <- c("Elite", "All-Star", "Starter", "Rotation", "Roster", "Out of League")

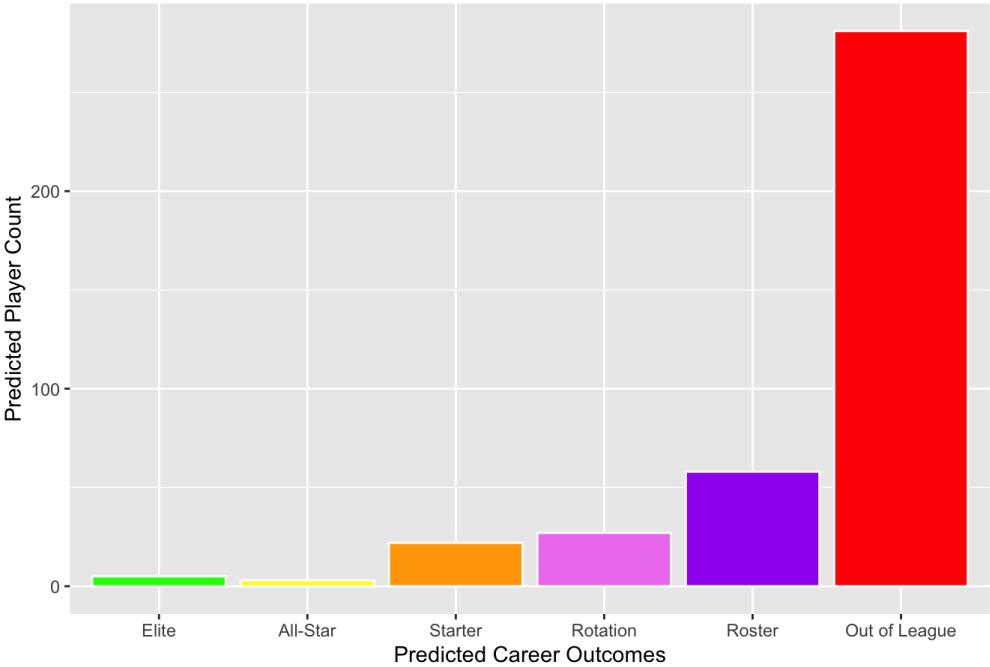
outcome_predictions$testing_predictions <- factor(outcome_predictions$testing_predictions, levels = desired_order)

custom_colors <- c("Elite" = "green",
                   "All-Star" = "yellow",
                   "Starter" = "orange",
                   "Rotation" = "violet",
                   "Roster" = "purple",
                   "Out of League" = "red")

ggplot(outcome_predictions, aes(x = testing_predictions)) +
  geom_bar(stat = "count", fill = custom_colors, color = 'white') +
  labs(title = "Predicted Outcome Counts for All 6 Classes",
       x = 'Predicted Career Outcomes', y = "Predicted Player Count")
```



Predicted Outcome Counts for All 6 Classes




```
###Importance of Variables in Random Forest Model###
variable_importance <- importance(rf_model)

variable_importance_df <- as.data.frame(variable_importance)

gini_threshold <- 10
important_vars <- variable_importance_df %>%
  filter(MeanDecreaseGini > gini_threshold)

ggplot(important_vars, aes(x = reorder(rownames(important_vars), MeanDecreaseGini), y = Mean
  geom_bar(stat = "identity", fill = "blue") +
  labs(title = "Feature Importance Plot",
    x = "Variable",
    y = "Mean Decrease in Gini Score") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Feature Importance Plot





```
stats_vs_outcomes_dataset <- left_join(total_predictions, testing_dataset_aggregated, by = "
View(stats_vs_outcomes_dataset)

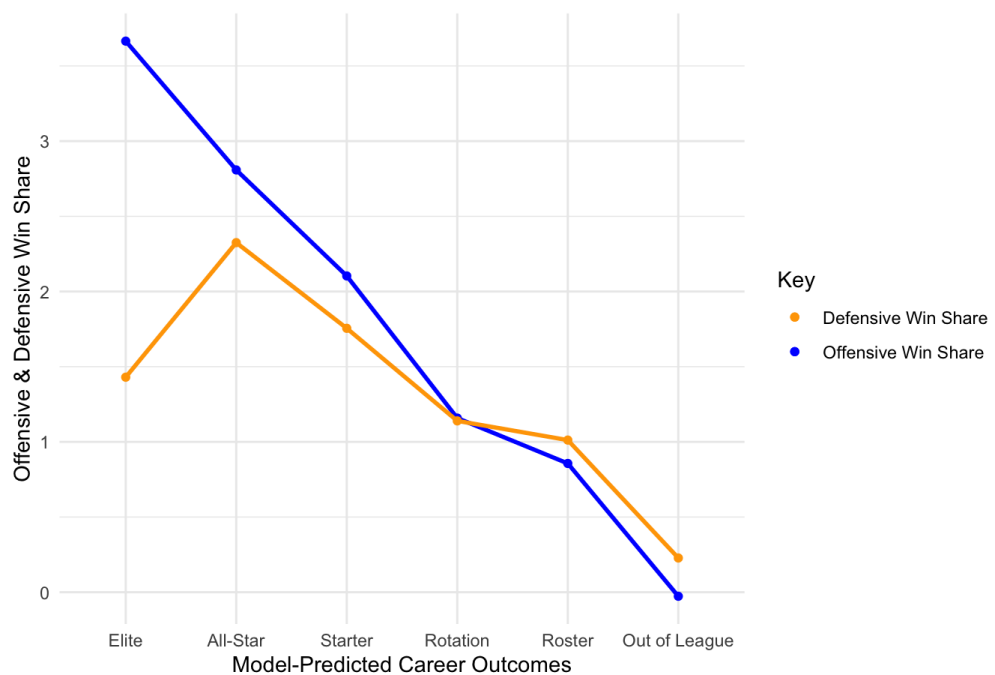
win_shares_by_outcome <- stats_vs_outcomes_dataset %>%
  group_by(`Model Predictions`) %>%
  summarize(Offensive_Win_Share_by_Outcome = mean(average_OWS),
            Defensive_Win_Share_by_Outcome = mean(average_DWS))
View(win_shares_by_outcome)

outcome_order <- c("Elite", "All-Star", "Starter", "Rotation", "Roster", "Out of League")

win_shares_by_outcome$`Model Predictions` <- factor(win_shares_by_outcome$`Model Predictions
rder)

ggplot(win_shares_by_outcome, aes(x = `Model Predictions`)) +
  geom_point(aes(y = Offensive_Win_Share_by_Outcome, color = "Offensive Win Share")) +
  geom_line(aes(y = Offensive_Win_Share_by_Outcome, group = 1), linewidth = 1, color = "blue") +
  geom_point(aes(y = Defensive_Win_Share_by_Outcome, color = "Defensive Win Share")) +
  geom_line(aes(y = Defensive_Win_Share_by_Outcome, group = 1), linewidth = 1, color = "orange") +
  labs(title = "Mean Win Share Trends Across Different Predicted Career Outcomes",
       x = "Model-Predicted Career Outcomes", y = "Offensive & Defensive Win Share") +
  scale_color_manual(name = "Key", values = c("Offensive Win Share" = "blue", "Defensive Win
+
  theme_minimal()
```

Mean Win Share Trends Across Different Predicted Career Outcomes



PREDICTIONS FOR SGA, ZION, WISEMAN, AND GIDDEY###

```
players_of_interest <- c("Shai Gilgeous-Alexander", "Zion Williamson", "James Wiseman", "Jos
predictions_specific_players <- outcome_predictions %>%
  filter(player %in% players_of_interest)

predictions_specific_players
```

##	player	testing_predictions
## 159	James Wiseman	Out of League
## 201	Josh Giddey	Roster
## 337	Shai Gilgeous-Alexander	All-Star
## 395	Zion Williamson	Elite



#If for any reason, the View() does not work, in order of SGA, Zion, Wiseman, and Giddey, my m as all-star, elite, out of league, and roster, respectively.

###BONUS: HTML TABLE WITH PREDICTIONS PLAYERS DRAFTED BETWEEN 2019-2021###

#I want to get all of the probabilities of each distinct outcome, as well as the predicted c r in one dataset. Important for next step#

```
total_predictions <- left_join(outcome_predictions, probability_predictions, by = "player")
```

#Now, I want to go back and merge this dataset that has the player name, probability of each ed outcome, with a previous dataset that also includes each players draft year#

##First, I'm going to make a copy of my master data set, filter for players drafted between st bring the variables 'player' and 'draftyear'##

```
HTML_player_dataset <- model_data %>%
  filter(draftyear >= 2018 & draftyear <= 2021) %>%
  select(player, draftyear)
```

#Now I merge the two datasets on their player so all 2018-2021 players will have player name ilities, and predictions. I will also group and filter so only one record shows per player, draft year.#

```
HTML_table_predictions <- left_join(total_predictions, HTML_player_dataset, by = 'player') %
  filter(draftyear >=2019) %>%
  group_by(player, draftyear) %>%
  filter(row_number() == 1) %>%
  ungroup()
```

```
View(HTML_table_predictions)
```

#I will install and load in reactable now#

```
library(reactable)
```

```
reactable(HTML_table_predictions)
```

player	testing_pred ictions	All-Star	Elite	Out of League	Roster	Rotation
Aaron Henry	Out of League	0	0	0.748	0.23	0.006
Aaron Nesmith	Out of League	0	0	0.648	0.172	0.158
Aaron Wiggins	Rotation	0	0	0.23	0.282	0.36
Adam Mokoka	Out of League	0	0	0.964	0.028	0.008
Ade Murkey	Out of League	0	0	1	0	0
Admiral Schofield	Out of League	0	0	0.93	0.044	0.022
Ahmad Caver	Out of League	0	0	0.908	0.082	0.006
Aleem Ford	Out of League	0	0	0.746	0.166	0.034



Aleksej Pokusevski	Out of League	0	0	0.526	0.23	0.188
Alen Smailagic	Out of League	0	0	0.844	0.052	0.104

1–10 of 294 rows

Previous123

```
#I didn't like the way the columns were named and want to make the table look a little bit more
entation table. I had to search for how to work with reactable but I believe I got it to wor

colnames(HTML_table_predictions)[1:9] <- c("Player", "Model-Based Predictions", "Elite", "Al
ue", "Roster", "Rotation", "Starter", "Draft Year")

Final_HTML_table <- HTML_table_predictions[, c("Player", "Model-Based Predictions", "Elite",
League", "Roster", "Rotation", "Starter", "Draft Year")]
Final_HTML_table$Elite <- round(Final_HTML_table$Elite, 2)
Final_HTML_table$'All-Star' <- round(Final_HTML_table$'All-Star', 2)
Final_HTML_table$Starter <- round(Final_HTML_table$Starter, 2)
Final_HTML_table$Rotation <- round(Final_HTML_table$Rotation, 2)
Final_HTML_table$Roster <- round(Final_HTML_table$Roster, 2)
Final_HTML_table$'Out of League' <- round(Final_HTML_table$'Out of League', 2)

reactable(Final_HTML_table,
  striped = TRUE,
  bordered = TRUE,
  highlight = TRUE,
  searchable = TRUE,
)
```

Player	Model-Based Predictions	Elite	All-Star	Out of League	Roster	Rotation
Aaron Henry	Out of League	0	0	0.75	0.23	0.01
Aaron Nesmith	Out of League	0	0	0.65	0.17	0.16
Aaron Wiggins	Rotation	0	0	0.23	0.28	0.36
Adam Mokoka	Out of League	0	0	0.96	0.03	0.01
Ade Murkey	Out of League	0	0	1	0	0
Admiral Schofield	Out of League	0	0	0.93	0.04	0.02
Ahmad Caver	Out of League	0	0	0.91	0.08	0.01
Aleem Ford	Out of League	0	0	0.75	0.17	0.03
Aleksej Pokusevski	Out of League	0	0	0.53	0.23	0.19
Alen Smailagic	Out of League	0	0	0.84	0.05	0.1



Part 2 – Predicting Team Stats

In this section, we're going to introduce a simple way to predict team offensive rebound percent in the next game and improve those predictions.

Question 1

Using the `rebounding_data` dataset, we'll predict a team's next game's offensive rebounding percent to be their average percent in all prior games. On a single game level, offensive rebounding percent is the number of offensive rebounds divided by the number of offensive rebound "chances" (essentially the team's missed shots). On a multi-game sample, it should be the total number of offensive rebounds divided by the total number of offensive rebound chances.

Please calculate what OKC's predicted offensive rebound percent is for game 81 in the data. That is, use games 1-80

#In summation, I took the team rebounding data, and filtered it only to show OKC's statistics and then used the summarize function to add up all offensive rebounds and all chances to get an offensive rebounding percent for OKC, and divided it to find OKC's offensive rebounding % in games 1-80, as well as its predicted value for game 81.

```
thunder_rebounding <- rebounding_data %>%
  filter(team == 'OKC') %>%
  summarise(
    total_off_rebounds = sum(offensive_rebounds),
    total_off_rebounds_chances = sum(off_rebound_chances),
    game_81_oreb_precent = total_off_rebounds / total_off_rebounds_chances
  )

thunder_rebounding
```

```
## # A tibble: 1 × 3
##   total_off_rebounds total_off_rebounds_chances game_81_oreb_precent
##           <dbl>           <dbl>           <dbl>
## 1           1226           4237           0.289
```

ANSWER 1:

28.9%

Question 2

There are a few limitations to the method we used above. For example, if a team has a great offensive rebounder who is injured this season but will be out due to an injury for the next game, we might reasonably predict a lower team offensive rebounding percent for the next game.

Please discuss how you would think about changing our original model to better account for missing players. You do not need to implement any changes, and you can assume you have access to any reasonable data that isn't provided in this problem. Be concise with your answer.

ANSWER 2: While I believe that using team statistics can be very helpful to analyze trends and find general insights into the performance of a team's offensive rebounding ability, especially over a longer sample size of a full season, it may not be the best way to predict a value for a single game if there are better measurables at hand. For example, within this situation, if a team's best rebounder is missing for the 81st game but being present in most of the team's previous games that contributed to their rebounding prediction, using individual player statistics to help add in a rough estimate of each player's individual offensive rebounding ability could help with the game-to-game variance in prediction. I was thinking it would be helpful to use each player's offensive rebounding statistics on a per game level. Such statistics could include total offensive rebounds, total offensive rebounding chances, their contribution to total offensive rebounds when they are on the floor, the amount of minutes they play, and the number of times they spend within five feet of the basket on offense as a measure of how often they are looking to rebound, among other statistics that could help build a model that analyzes these measures and determines each player's level of contribution to the team's offensive rebounding.

Question 3

In question 2, you saw and discussed how to deal with one weakness of the model. For this question, please write about the other weaknesses of the simple average model you made in question 1 and discuss how you would deal with each of them. You can discuss one weakness and discuss how you'd fix that weakness, then move onto the next issue, or you can start by explaining multiple weaknesses and how to deal with them.



original approach and discuss one overall modeling methodology you'd use that gets around most or all of them. Agree on any code or implement any changes, and you can assume you have access to any reasonable data that isn't provided. Your answer should be clear and concise with your answer.

ANSWER 3: Using averages of team statistics to predict the next game's value is difficult as there is little to no game-by-game changes to statistical output in an NBA season, as shown by the situation in Question 2 that a player is injured. It is also worth noting that who you are playing is very important in individual offensive rebounding statistics. Teams have either more players with skills and physical attributes that gives them advantages on the glass, probably leading to higher team defensive rebounding numbers. The opponent may also have a coach that emphasizes a play style that creates more opportunities based on their defensive positioning, or a combination of both skill and play style. This greatly affects a team's offensive rebounding game-by-game, and controlling a possible regression or model to have values for offensive and defensive rebounding ability is imperative. This includes total defensive rebounding percentage and individual player defensive rebounding percentage. How valuable their individual players playing in the next game are to the team's overall defensive rebounding is also a factor. The most efficient way to find a league average of total defensive rebounding based on defensive rebounding percentage and set that value equal to 1. Then if you are playing against a certain team and are aware of the availability of individual player statistics to predict a value of defensive rebounding percentage will give you a better look at defensive rebounds and percent should be, as compared to what you would have found if you just added together all of the games defensive rebounding numbers and chances and found an average. This new number, driven by individual player statistics, to what I thought would make sense to use in question 2 in regards to your own team's offensive rebounding percentage compared to the league average. Dividing the number found by the league average tells you in a percentage how much the opponent is than the league average. If the number is greater than 1, say 1.1, that means the opponent is 10% better than the league average. If the number is less than 1, say 0.9, that means the opponent is 10% worse than the league average. This would normally be suspect, so you may adjust our average offensive rebounding percentage to reflect playing against a team that is better or worse than the league average. One way to go about inputting how important different opponents are on a game-by-game level to the variance in offensive rebounding and how to possibly correct this in our basic dataset so as to alter our team's predicted values as well. A second model we are predicting our values. There are many other ways to predict and evaluate certain values, and using simple averages in the context of an NBA game is not ideal. For example, take using a multivariate regression analysis that incorporates all of the data necessary to incorporate in our model. This could be individual player statistics and their role in offensive rebounding, both for our team's offensive rebounding and for opponents defensive rebounding. Adding certain game level factors such as home versus away, how many days rest prior to the game a team had, what is a team's pace, how many different shots do they take per game (three pointers, mid ranges, free throws, and layups/dunks all are not created equal in terms of offensive rebounding probability). All of these factors are super important to have in a larger offensive rebounding model. This model potentially use in a regression model that predicts a team's offensive rebounding using all of these factors. The coefficients that essentially display how much value each factor plays into total offensive rebounding per game. This model allows you to plug in all the necessary values that change game by game to predict the next game's offensive rebounding precisely. If you know what team you are playing, where you are playing, if any specific player is unavailable or injured, the prior rebounding statistics and shot statistics, the model will show a more accurate representation of offensive rebounding for that game as compared to the method we used in the first question.